kNN TigerFish Report

Dane Acena

Problem Description

Given a data set from the Clemson Wildlife and Fisheries Biology graduate students of two newly discovered species of fish in Lake Hartwell, develop a kNN classification algorithm that will determine a fish is a TigerFish1 (the positive case) or a TigerFish0.

Data Description

The initial data consisted of 300 records representing features of either TigerFish1 species or TigerFish0 species with three tab-delimited entries. The first two are floats indicating the measured body length and dorsal fin length of each fish, respectively. The last element is a digit, either "1" or "0" identifying the species of fish as either "TigerFish1" or "TigerFish0". A plot of the initial dataset is as shown in Figure 1.

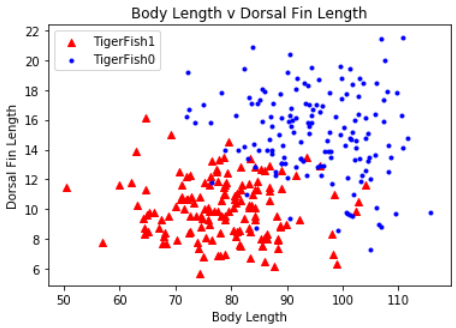

Figure 1: The Initial Data Set

Training a kNN Algorithm

The k Nearest Neighbor algorithm was developed using 5-fold cross-validation. When the data set is loaded it gets randomized and split 80-20. The 80% goes to a *Training Set,* and the remaining 20% goes to a *Test set*, which will be set aside be put to use later on. The algorithm further divided *Training Set* into five folds with 48 records each. From those five sets, the algorithm recursively creates a *Train[i].txt* (contains 192 records each) combining four sets and create *Val[i].txt* (contains 48 records each) with the leftover fold. **validationMode()** recursively executes Train[i].txt and Val[i].txt via **crossValidate()** with odd values of k from 1 through 21. In each iteration of the test, each record of the Val[i].txt set gets tested against the Train[i].txt set and **getNeighbors()** looks at the *k* points nearest the *val* record using **euclideanDistance()**. Then, **getResponse()** determines the majority type of neighbor**.** The accuracy and error are acquired through **getAccuracy()** and **getError()** respectively.

For each test iteration of each test set and k value, the number misclassification is recorded in a tab-delimited file, as shown in Figure 2.

k	1	3	5	7	9	11	13	15	17	19	21
Test1 Errors	6	5	5	6	6	6	6	7	6	6	6
Test2 Errors	3	3	3	3	4	4	5	5	5	5	5
Test3 Errors	9	7	7	6	7	7	7	7	7	7	7
Test4 Errors	5	6	5	3	4	4	4	4	5	4	5
Test5 Errors	5	2	2	2	2	2	2	2	2	1	1
TOTALS	**28**	**23**	**22**	**20**	**23**	**23**	**24**	**25**	**25**	**23**	**24**

Figure 2: Misclassifications for different values of k on the five training sets

Additionally, to visually illustrate the accuracy a plot was also produced by the algorithm, as shown in Figure 3.

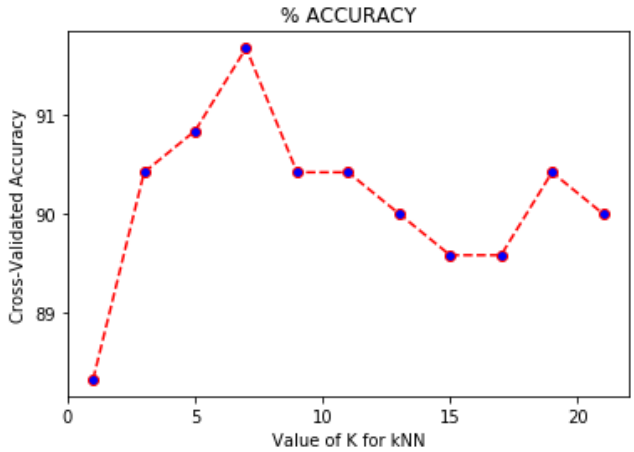

Figure 3: Average accuracy for different values of k (PLACEHOLDER)

As shown above, k=7 provided the best accuracy. K=7 was chosen for the test set with the k Nearest Neighbor algorithm.

Results

Results of the validation of the test set against the training set are shown in the confusion matrix as shown in figure 4

The test set contained 34 records representing TigerFish0 and 26 records representing TigerFish1. 57 of the 60 records identified correctly; denoting an *accuracy* of 0.95. *Precision* is 0.96, out of the 25 that were predicted to be TigerFish1, one was TigerFish0. The *recall* is 0.92, two out of the 35 that were predicted to be TigerFish0 are misclassified. The **F1** score is 0.94

		Predicted TigerFish1	
		N	Y
Actual TigerFish1	N	TN =33	FP =1
	Y	FN =2	TP =24

Figure 4: Confusion Matrix